

Beyond the Chatbox

Egor Kotov

Contents

	Introduction	7
	Overview	7
	Workshop Details	7
	Before the Workshop: Accounts & Setup	8
	Recommended Accounts to Register	8
	Coding Agents We Will Use (and Why)	8
1	Introduction to LLM Agents	11
1.1	What is an LLM Agent?	11
1.1.1	Key Characteristics	11
1.2	Why Use Agents for Research?	11
2	Environment Setup	13
2.1	 Launching the Environment	13
2.1.1	 Auto-Pause & Resume	14
2.2	 Managing Your Session & Quotas	14
2.2.1	Monthly Free Quotas (2026)	14
2.2.2	Checking Remaining Quota	14
2.2.3	Stopping Your Session Manually	15
3	Navigating the Workshop Environment	17
3.1	The VS Code Interface	17
3.2	Using folders as projects	17
3.2.1	Open project in a new window	17
3.2.2	Open terminal in a specific folder	17
3.2.3	Launching an Agent	18
3.3	Global file explorer	19
4	Accessing models and authenticating agents	21
4.1	Authenticating Agents	21
4.2	OpenCode	21
4.2.1	OpenCode Zen	22
4.2.2	OpenRouter	22
4.2.3	NVIDIA NIM/Build	22
4.3	Google	22
4.3.1	Google Antigravity CLI	22
4.3.2	Google Gemini CLI	23

5	The Core Brain: Models and Prompts	27
5.1	Models	27
5.2	Model vs. Agent System Prompts	27
5.2.1	The Model System Prompt	27
5.2.2	The Agent System Prompt	27
5.3	Configuration Files: Global and Local Instructions	27
5.3.1	Global vs. Local Instructions	27
5.4	Prompt Layering Architecture	28
6	Task 1: Exploring System Prompts	31
6.1	Mini-Exercise 1: Extracting the Agent Prompt	31
6.2	Mini-Exercise 2: Dumping Instructions to a File	31
6.3	Mini-Exercise 3: Changing Personas	31
6.4	Mini-exercise 4: Guiding the Agent to do things in a certain way	32

II

Safety

7	Safety and Isolation: Sandboxes and Permissions	35
7.1	Sanboxing and Containerization	35
7.2	Ignore files	35
7.3	Permission Systems	36
7.4	Bottom line	36
8	Task 2: Boundaries and Permissions	37

III

Capabilities

9	Capabilities: Tools and Skills	41
9.1	Tools: The Agent's Hands	41
9.2	Skills: Specialized Capabilities	41
10	Task 2: Working with Tools and Skills	43
10.1	Setting the Scene	43
10.2	Step 1: Exploratory Analysis	43
10.3	Step 2: Activating a Skill	43
10.4	Step 3: Fixing a Bug	43

IV

State Management

11	State Management: Context and Memory	47
11.1	The Context Window	47
11.2	Memory Systems	47
11.3	Agentic vs. Naive Approaches	47

12	Task 3: The Context Limit Challenge	49
12.1	Step 1: Generating the “Mega-File”	49
12.2	Step 2: The Naive Reality Check	49
12.3	Step 3: The Agentic Solution	49
12.4	Why this matters	49
	References	51
	Production tools	51
	Website libraries	51
	Local project code	52
	Content license	52

Introduction

⚠ Warning

THIS IS STILL WORK IN PROGRESS - EXPECT BROKEN LINKS AND MISSING CONTENT



Figure 0.1: © Image generated by Google Gemini | Nanobanana

Overview

This interactive tutorial and workshop provides a practical introduction to using large language models (LLMs) in coding agents, featuring applied examples in population research. It is organized by [Egor Kotov](#) with support from the [Max Planck Institute for Demographic Research](#).

Workshop Details

- **Date:** June 03, 2026
- **Time:** 13:30-16:30
- **Location:** Complesso Belmeloro, room “Lab O”, [Via Beniamino Andreatta, 8, 40126 Bologna BO, Italy](#)

Abstract

The software industry has been enjoying the benefits of coding agents for over a year now, with even top software engineers now having AI assistants generate a vast majority of their code. Now, the tools for agentic coding are mature enough for anyone to try.

Unlike many educational materials in this field, this workshop will also address the new vectors of attack that novice users might expose themselves to by using coding agents blindly. We will cover key security risks associated with agentic systems, including data leakage, unsafe code execution, and unintended modifications to analysis code and research materials. To mitigate these risks, we will introduce accessible strategies such as containerization, sandboxing, and workflow isolation.

Participants will also see how the agents themselves can assist in setting up and using these protective tools to improve transparency and reproducibility. The workshop will combine short demonstrations with applied examples relevant to population research. Participants are encouraged to bring their own projects or ideas, which we will use to explore how LLM agents can be integrated into real research workflows. For those without a project, guided exercises and example tasks will be provided.

Before the Workshop: Accounts & Setup

To make the most of our hands-on sessions, please complete the following steps before the workshop on **June 3, 2026**.

! Data Privacy & Security Warning

Free tiers for these LLM services typically **use your prompts and uploaded code for model training. Do not use these tools on private, sensitive, or proprietary research data!** Only use public or non-sensitive datasets during the workshop and for general testing.

Recommended Accounts to Register

Please sign up for the following services in advance to ensure you have enough free model usage quota during the exercises:

- **Google Account** (A personal/secondary account is perfect) — Used for **Gemini CLI** and **Antigravity CLI**. No credit card is required.
- **OpenCode Zen** (opencode.ai/zen) — Gives you free model access via the **OpenCode** agent. You can view the list of available models and details on [OpenCode Zen pricing/documentation](#).
- **OpenRouter** (openrouter.ai) — Allows accessing a wide variety of models. You'll need to generate a free [API Key](#) after registering, which can be used to access their catalog of [free models](#).
- *Optional:* **NVIDIA NIM/Build** (build.nvidia.com) — Offers [free preview models](#) (note: account activation requires SMS verification via a mobile phone number).

i Note

If you have a subscription to another LLM service (such as OpenAI ChatGPT Plus, Claude Pro, Gemini Advanced, or Mistral) and prefer to use it during the workshop, please feel free to do so. However, our guided exercises will focus on the tools listed above to ensure everyone has access to the same models and features.

The skills you will learn here are highly transferable to other AI coding assistants, though setup steps and specific features may vary by provider. Please refer to the documentation of your chosen agent for setup instructions, and don't hesitate to ask the agent itself how to configure it for your desired models and features.

⚠ API & Account Safety Warning

Do **NOT** attempt to authenticate your Google or Anthropic (Claude) accounts directly within OpenCode. Their terms of service prohibit use of quota based subscriptions in third-party clients, and doing so can result in your accounts being permanently banned. OpenAI is currently a bit more flexible on this at the time of writing, but always verify provider policies.

Coding Agents We Will Use (and Why)

The coding agents are already pre-installed and configured inside the workshop's Docker container / Codespaces environment. We will focus on:

- **Google Antigravity CLI** (and its predecessor **Gemini CLI**)
 - *Why:* It provides a free daily quota to test cutting-edge models (like [Gemini 3.5 Flash](#), [Gemini 3.1 Pro](#), and [Claude](#) models) directly from your Google account.

- **OpenCode CLI**

- *Why:* A flexible, open-source-friendly agent that can connect to multiple backends. We will use it with free models from **OpenCode Zen**, **OpenRouter**, and **NVIDIA NIM** to show how easy it is to switch between different model providers without vendor lock-in.

i Note

What is **CLI**? It stands for “Command Line Interface” and means that you will interact with the agent by typing commands in a terminal (similar to your R console), rather than through a graphical user interface (GUI). This allows for more flexibility and control over the agent’s behavior, as well as easier integration into coding workflows. This also gives you better insight into how the agent works under the hood, which is important for understanding its capabilities and limitations. You are of course free to use the same agents through their graphical interfaces too. We are also limited by the current tutorial environment and tools that are free to use. In our tutorial environment Gemini and Antigravity are only available as CLI, OpenCode is also CLI, but you will learn how to run it in graphical mode as well.

To get started with the workshop materials, please navigate through the sections in the navigation bar on the left or with pagination links in the bottom.

1. Introduction to LLM Agents

This section covers the basics of what LLM agents are and how they differ from standard chat interfaces.

1.1 What is an LLM Agent?

Unlike a standard LLM chat (like OpenAI ChatGPT or Google Gemini, or Claude web interface), an **agent** is a system that uses an LLM as its “reasoning engine” to perform tasks by interacting with the world. Agent essentially adds hands to an LLM, allowing it to execute code, use tools, and access external information.

Most LLM chat interfaces today already have some tools (e.g. they can write and execute (mostly Python) code, or access a search engine). They can also be connected to external tools or data via plugins or APIs. However, they are still primarily designed for conversational interactions and may not be optimized for complex task execution or multi-step reasoning. Though there is nothing fundamentally preventing them from doing that in the future.

1.1.1 Key Characteristics

- **Tool Use:** Agents can use tools such as search engines, calculators, or terminal commands.
- **Autonomy:** They can break down a complex goal into smaller steps and execute them sequentially.
- **Feedback Loop:** They can observe the output of a tool (e.g., a code error, a statistics model output, a plot) and adjust their next action accordingly.

1.2 Why Use Agents for Research?

In academic research, agents can:

- **Automate Data Cleaning:** Write and run scripts to fix inconsistencies in large datasets.
- **Literature Review:** Search, summarize, and synthesize findings from hundreds of papers.
- **Reproducible Analysis:** Generate documented code that follows best practices.

2. Environment Setup

To ensure a consistent, secure, and hassle-free experience, we will use a pre-configured cloud-based development environment. We will use **GitHub Codespaces**.

💡 ⌚ Start the Setup Early and Focus on the Presentations!

First-time environment setup and container building can take **up to 10 minutes** to compile and prepare.

We highly recommend launching the codespace right at the beginning of the workshop. While the setup runs automatically in the background, you can continue to fully focus on the presentations, live demonstrations, and slides shown by the tutor.

2.1 🛠️ Launching the Environment

To launch the Codespaces environment manually from the GitHub interface, follow these steps:

1. Navigate to the [workshop repository](#).
2. Click the green **Code** button in the top right.
3. Select the **Codespaces** tab.
4. Click **Create codespace on main** (or click the **...** menu to configure machine size).




Figure 2.2: Start GitHub Codespaces

Free & Secure Environment

- **It is 100% Free:** GitHub provides a generous free tier of **60 hours** of Codespaces usage per month for every personal account. No credit card, billing info, or payment is required.
- **It is Completely Secure:** The environment runs inside a secure, isolated container (virtual machine) in the cloud. It does not touch your physical computer, install any local files, or run any software on your machine. It belongs to you and it is also mostly safe to enter your account data (e.g. Google account, especially if secondary or test account, API keys, other accounts for subscription services or services we will use during the workshop) without worrying about security or privacy issues.

[media/videos/start-codespace.mp4](#)

Figure 2.3: Codespace startup process (sped up)

2.1.1 Auto-Pause & Resume

To save your free hours, the environment automatically pauses after 30 minutes of inactivity. When you restart it, your files, open tabs, and terminal history are fully preserved—just like waking up a laptop from sleep (though you may need to restart any active terminal commands or running applications).

2.2 Managing Your Session & Quotas

GitHub Codespaces provides a generous free tier, but it is important to understand how usage is consumed to avoid running out of credits.

2.2.1 Monthly Free Quotas (2026)

The following quotas apply to personal accounts.¹

Account Type	Actual Usage Hours	Storage (per month)
GitHub Free	60 hours	15 GB
GitHub Pro	90 hours	20 GB

Storage Persistence Quota

Storage (15GB/20GB) is consumed even when the codespace is **stopped**. To save your storage quota, make sure to **delete** codespaces you no longer need at github.com/codespaces once the workshop is fully over.

2.2.2 Checking Remaining Quota

Tip

You can monitor your real-time usage and remaining balance at any time:

1. Go to your [GitHub Settings > Billing and plans](#).
2. Look for the **Codespaces** section under **Usage this month** or try this direct link: [GitHub Codespaces Usage](#).

¹While the [official billing documentation](#) lists these as 120 and 180 **core hours**, the minimum instance size is 2 CPU cores. This means every 1 hour of real-time work consumes 2 core hours, resulting in the actual limits shown below.

2.2.3 Stopping Your Session Manually

When you are done with a session, you can stop the codespace manually to instantly stop consuming usage hours:




Figure 2.4: Stop GitHub Codespace

3. Navigating the Workshop Environment

Before we start interacting with agents, let's take a moment to understand the “cockpit” of our workshop: the VS Code / GitHub Codespaces environment.

3.1 The VS Code Interface

The workshop runs in a specialized container that includes all necessary tools. See the [figure](#) below for the main elements of the VS Code interface in the Codespace environment. Click the image to zoom in.




Figure 3.5: Main elements of VS Code interface.

3.2 Using folders as projects

3.2.1 Open project in a new window

In VS Code a folder is considered a project or workspace.

To open any folder as a new project/workspace, right-click on the folder and choose **Open New Workbench Here**.

[media/videos/vscode-open-folder-as-project-new-window.mp4](#)

Figure 3.6: Opening a folder as a new project in VS Code.

This allows to easily manage this workspace as well as Git repository for this specific folder.

3.2.2 Open terminal in a specific folder

For quick agent tasks in any folder without opening a new window, you can also open a terminal directly in that folder. Right-click on the folder and select **Open in Integrated Terminal**.

[media/videos/vscode-open-folder-in-terminal.mp4](#)

Figure 3.7: Opening terminal in a specific folder in VS Code.

If you open the folder in terminal this way and start an agent there, it will treat this directory as its workspace and will have access to all files in this directory and its subdirectories. But you will not be able to just as easily manage this folder as a separate project in VS Code or Git, as the whole Codespace will be treated as one project by VS Code.

Note

If at any point you feel lost, you can always reopen the default folder that the Codespaces started with by clicking the burger menu in the top left corner of VS Code and selecting **File -> Open Folder** and pasting `/workspace/` there. Or use the **File -> Open Recent** menu to find the folder you want to open.




Figure 3.8: Opening a folder in VS Code.

3.2.3 Launching an Agent

In a terminal in the bottom of VS Code you can run any terminal command, you can run R or Python consoles, and also agents.

Commands to launch the agents will be provided in the next sections.

If you are not familiar with terminal, common Linux commands are:

- `ls` - list files in the current directory
- `cd` - change directory (e.g. `cd folder_name` to go into a folder, `cd ..` to go back to the parent directory)
- `pwd` - print working directory, shows the current directory you are in
- `mkdir` - make a new directory (e.g. `mkdir new_folder` to create a new folder)
- `rm` - remove a file or folder (e.g. `rm file_name` to remove a file, `rm -r folder_name` to remove a folder and all its contents). Just make sure to never run `rm -rf /` command, as it will delete everything in the container and break your Codespace.

You don't need to memorize these, as agent can navigate the terminal and file system for you. All you need to do is give it instructions in natural language and allow it to do things for you.

3.3 Global file explorer

VS Code is designed to use specific folders as projects or workspaces, so by default you only see the files in the current directory in the file explorer pane.

If you need to browse all files in the Codespace, you can use **File Explorer** extension preinstalled for you in the Codespace:

[media/videos/vscode-file-explorer.mp4](#)

Figure 3.9: Using **File Explorer** extension to browse all files in the Codespace.

4. Accessing models and authenticating agents

4.1 Authenticating Agents

I suggest you start with more basic models available for free with OpenCode Zen, OpenRouter and optionally NVIDIA NIM/Build, and then try the more powerful frontier models available in Gemini CLI and Antigravity CLI.

Authenticate at least one of the services below before proceeding. Saving authentication credentials in the Codespaces environment is safe, as it runs in an isolated container in the cloud and only you have access to it via your GitHub account. However, do make sure to use expiration date when setting up API keys or logout from the Google account (for Gemini models) after the workshop to avoid any potential security issues in the future. Also remember to [delete the Codespace](#) once the workshop is fully over.

4.2 OpenCode

You can just start using the OpenCode agent immediately after launch without API key, but model quota will be small, as you are likely sharing an IP address with several other workshop participants.

Just start OpenCode in a new terminal with command:

```
opencode
```




Figure 4.10: Opening a new terminal in VS Code.

Once it starts, use `/model` slash-command to select a model and type `free` to see all available free models.

4.2.1 OpenCode Zen

OpenCode Zen (opencode.ai/zen) gives you free model access via the **OpenCode** agent. You can view the list of available models and details on [OpenCode Zen pricing/documentation](#). Get a free account, generate API key and add it in OpenCode.

[media/videos/opencode-zen-api.mp4](#)

Figure 4.11: Adding OpenCode Zen API key to OpenCode agent.

4.2.2 OpenRouter

OpenRouter (openrouter.ai) allows accessing a wide variety of models. You'll need to generate a free **API Key** after registering, which can be used to access their catalog of [free models](#).

[media/videos/opencode-openrouter-api.mp4](#)

Figure 4.12: Adding OpenRouter API key to OpenCode agent.

4.2.3 NVIDIA NIM/Build

This is completely optional, as it takes a bit more time to setup and requires SMS verification. You are probably going to be fine with OpenCode Zen and OpenRouter free models + models in Google's Gemini CLI and Antigravity CLI.

NVIDIA NIM/Build (build.nvidia.com) offers [free preview models](#).

[media/videos/opencode-nvidia-api.mp4](#)

Figure 4.13: Adding NVIDIA NIM API key to OpenCode agent.

💡 OpenCode Graphical user interface

If you want to try OpenCode's graphical user interface instead of the command line, you can launch it with `opencode web` command in the terminal or use **OpenChamber VS Code** extension.

[media/videos/opencode-gui.mp4](#)

Figure 4.14: Running OpenCode's graphical user interface with `opencode web` command or OpenChamber VS Code extension.

4.3 Google

Google has two separate agents with different model access and quotas: **Gemini CLI** and **Antigravity CLI**. Gemini CLI models will be disabled for most users mid-June 2026, but you can still use it to try out the models for free. Going forward, Antigravity CLI will be the main way to access Google's models, and it also provides a free daily quota.

4.3.1 Google Antigravity CLI

The **Google Antigravity CLI** (and its predecessor **Gemini CLI**) provides a free daily quota to test cutting-edge models (like **Gemini 3.5 Flash**, **Gemini 3.1 Pro**, and **Claude** models). You can authenticate it using your Google account credentials, no credit card required.

To start Antigravity CLI, in new terminal run:

```
agy
```

[media/videos/agy-setup.mp4](#)

Figure 4.15: Authenticating Google Antigravity CLI with your Google account.

4.3.2 Google Gemini CLI

The **Google Gemini CLI** is the predecessor to Antigravity CLI and provides free quota to **Gemini 3 Flash**, **Gemini 3.1 Pro** and a few other models. You can authenticate it using your Google account credentials, no credit card required.

To start Gemini CLI, in new terminal run:

```
gemini
```

[media/videos/gemini-cli-setup.mp4](#)

Figure 4.16: Authenticating Google Gemini CLI with your Google account.



The Brain

5	The Core Brain: Models and Prompts	27
5.1	Models	27
5.2	Model vs. Agent System Prompts	27
5.3	Configuration Files: Global and Local Instructions	27
5.4	Prompt Layering Architecture	28
6	Task 1: Exploring System Prompts .	31
6.1	Mini-Exercise 1: Extracting the Agent Prompt .	31
6.2	Mini-Exercise 2: Dumping Instructions to a File .	31
6.3	Mini-Exercise 3: Changing Personas	31
6.4	Mini-exercise 4: Guiding the Agent to do things in a certain way	32

5. The Core Brain: Models and Prompts

The most fundamental part of any coding agent is the **Model**—the “brain” that processes your requests. However, the model doesn’t work in a vacuum; it is wrapped in an **Agent Harness** (e.g. OpenCode, Claude Code, Codex, Google Antigravity CLI, many others).

5.1 Models

In this workshop we are going to use a variety of models, at the link below you can see a comparison of the characteristics, performance, and costs of the models we will be using, as well as some other popular models that you can try on your own.

[See models comparison](#)

5.2 Model vs. Agent System Prompts

An agent’s behavior is governed by layers of instructions. It is helpful to distinguish between the instructions given to the model and those that define the agent harness.

5.2.1 The Model System Prompt

This is the base set of rules provided by the model provider (e.g., OpenAI, Anthropic, Google). It typically contains safety guidelines and a persona like “*You are a helpful AI assistant.*” These instructions are often “hard-coded” or pre-configured by the provider before the agent harness adds its own layer.

5.2.2 The Agent System Prompt

The **Agent Harness** (the software you use to interact with the model) provides its own set of instructions. Depending on the harness, these might include:

- What it is (e.g. “*you are OpenCode, an agent that can read and write files, run code, and use tools to help users with coding tasks*”).
- How to use tools (like reading files or running bash commands).
- Where and how to access documentation on how to configure the agent harness (that is, you can actually ask most agent harnesses to change certain settings by asking them using natural language.)

! Important

Key Insight: The agent harness often does not have access to the underlying Model System Prompt. They are separate layers of instruction that are combined (concatenated) before being sent to the LLM provider that does the actual processing.

5.3 Configuration Files: Global and Local Instructions

The behavior of an agent harness is often guided by specific configuration files that provide context about you and your project. In our workshop, we focus on a hierarchy of **Instruction Files** (typically named `GEMINI.md` for Google’s models, `CLAUDE.md` for Anthropic’s models, and `AGENTS.md` for most other agents including OpenAI Codex).

5.3.1 Global vs. Local Instructions

Harnesses typically look for instructions at two levels:

1. **Global Instructions** (`~/.gemini/GEMINI.md` or `~/.agents/AGENTS.md`): These are your personal preferences that follow you across all projects. They might include your preferred coding style (e.g., “I always use the Tidyverse”), your language preferences, or general rules about how you want the agent to communicate.
2. **Local (Project) Instructions** (`./GEMINI.md` or `./AGENTS.md`): These are specific to a single repository. They define the overall project architecture, naming conventions, and project-specific workflows. Local instructions usually take precedence over global ones for that specific project in case of a conflict.

This may differ from agent to agent, so refer to the documentation of the specific agent harness ([OpenCode](#), [Antigravity CLI](#), [Gemini CLI](#), [Claude Code](#), [Codex](#)), or just as the agent itself.

5.4 Prompt Layering Architecture

The final prompt sent to the LLM service is a “sandwich” of different instruction layers. Understanding this hierarchy helps you write better prompts and understand agent behavior.




Figure 5.17: Prompt Layering Architecture: The hierarchy of instruction layers.

6. Task 1: Exploring System Prompts

In this task, you will interact with the **OpenCode** agent to understand the layers of instructions that govern its behavior.

Open the [workshop Codespace environment](#) if you haven't already.

Open the terminal in VS Code and launch the OpenCode agent with command:

```
opencode
```

Reminder how to get around the interface is in [Navigating the Interface](#).

Make sure you have selected a free model before you begin. See [Accessing models and authenticating agents](#) for instructions on how to do this.

6.1 Mini-Exercise 1: Extracting the Agent Prompt

Prompt to try:

```
"What is your agent system prompt? Can you also tell me what the underlying model system prompt is?"
```

Observation:

- Did the agent describe its system prompt?
- Did it know about its "Model" instructions?

 **Tip**

Try doing the same with Google Gemini CLI or Antigravity CLI and see if the results are different. Alternatively, ask the coding agent you might be using regularly.

6.2 Mini-Exercise 2: Dumping Instructions to a File

Now, let's test the agent's ability to interact with your workspace while explaining its rules.

Prompt to try:

```
"Summarize your core behavior rules and dump them into a file called `agent_instructions.txt` in this folder."
```

Action:

- Look at your file pane in VS Code. Did `agent_instructions.txt` appear?
- Open the file and see what the agent thinks its most important rules are.

6.3 Mini-Exercise 3: Changing Personas

Try giving the agent a specific "hat" to wear.

Prompt to try:

```
"From now on, act as a rockstar data scientist who prioritizes ultra-clean, vectorized R code. How would you explain the difference between a `for` loop and `lapply`?"
```

Observation:

- How did the tone and technical depth of the response change?

- This illustrates how your “User Instruction” layer sits on top of the fixed System Prompts.

6.4 Mini-exercise 4: Guiding the Agent to do things in a certain way

Here you will try to instruct the agent to do something in a specific way in your project directory to overcome some of its assumptions.

Prompt the model:

```
Use mtcars dataset (or any other built-in R dataset) to demonstrate to me how to do tidy statistical modelling using `tidymodels` R package. Create a reproducible quarto markdown report for me to run through interactively.
```

Then try this (in case the agent did not try to install the `tidymodels` package on its own):

```
install any packages needed to render the report and render it into html
```

At this point you might be asked to allow certain actions/tools the agent will need to do to complete the task, such as installing packages or writing files. In current environment, feel free to allow these actions and see how the agent tries to complete the task.

Notice how agent tries to install packages. Very likely it will use its training data to explicitly run `install.packages('tidymodels', repos = 'https://cloud.r-project.org/')` or similar. The problem is that in the Codespace you are running an Ubuntu Linux operating system, and on Linux R packages have to be built from source, which is very often very slow. The container you are in is perfectly setup to install packages from [Posit Package Manager](#) that provides pre-build binaries of R packages, which allows the installation to be as fast as on macOS or Windows. But the model chose to ignore this or did not bother to check.

So this is where you can press `Esc` key to interrupt the agent and give it some help. E.g. you can tell it to use `https://p3m.dev/cran/__linux__/noble/latest` for package installation, or you can create an `AGENTS.md` file in project root explaining briefly that it runs in Ubuntu Linux and packages should be installed from `https://p3m.dev/cran/__linux__/noble/latest` and reload the agent and try starting the conversation over. At this point the agent will already know how to do it properly in this specific environment.

Once you are done, close the agent by typing `exit` or `quit`.

II

Safety

7	Safety and Isolation: Sandboxes and Permissions	35
7.1	Sanboxing and Containerization	35
7.2	Ignore files	35
7.3	Permission Systems	36
7.4	Bottom line	36
8	Task 2: Boundaries and Permissions .	37

7. Safety and Isolation: Sandboxes and Permissions

Coding agents are incredibly powerful—they can delete files, run commands, and access the internet. This power comes with significant security risks. To mitigate these risks, agents operate within **Sandboxes** and under strict **Permission** systems.

7.1 Sanboxing and Containerization

A **Sandbox** is an isolated environment where the agent can work without affecting your main computer. Usually, you would like the agent to have access to its own folders (with system prompt, settings and skills) and to your specific project folder, but not anywhere else on your computer. This way, if the agent makes a mistake or runs a malicious command, it only affects the sandbox and possible files in your project, but not your entire system.

In this workshop, your entire environment is inside a Docker container. If the agent makes a mistake or runs a malicious command, it only affects the container, not your laptop.

All operating systems provide different ways of sandboxing or containerization, and many agent harnesses have built-in sandboxing features (but they still depend on your operating system). For example, there is well written concise documentation for [Gemini CLI](#) on the ways [it supports sandboxing](#), also see [Claude Code sandboxing](#), [Codex Sandboxing](#). For other agents, you might need a more custom setup, for example using [Docker Sandboxes](#) or [Virtual Machines](#) (e.g. [VirtualBox](#)) to create a isolated environment for the agent to run in and only connect your project folder to that environment. For most of these tools you have to have administrative access to your computer to set up the sandbox, there for for your work computer, ask your IT team if they can set up a sandboxed environment for you to run agents in.

Sandboxing also allows to allow the agent to be more autonomous and possibly also run in the so called 'YOLO' ('you only look once') or 'dangerously skip permissions' mode, where the agent essentially can do whatever it decides to do without asking for permissions. That allows you to stop babysitting the agent and let it do its work without interruptions, but it also means that if the agent makes a mistake or is tricked into doing something malicious, it can cause a lot of damage. So be very careful with this mode.

In this workshop in the provided Codespace Docker container you can feel, free to run the agents in this mode.

To run agents in 'YOLO' mode, you can either start them in this mode:

```
gemini -y
agy --dangerously-skip-permissions
# opencode does not have this mode yet, but it can be configured to just skip all
permissions in its settings file, ask the agent how to do it
```

Gemini CLI also allows swithing it on while running with **Ctrl+Y** hotkey.

7.2 Ignore files

When working on your tasks, agents often scan and read full files or parts of files in your projects. This data is of course then sent over to the model provider along with your prompt and chat history. This data may (and with free models we use in the workshop will) be used in future model and agent training. Therefore, you might want to exclude certain files from being read by the agent.

You can instruct the agent to ignore certain files with [.gitignore](#) file in the project root, or with a custom [AGENTS.md](#) file with instructions for the agent (e.g. "never read [.env](#) files"). The exact way to do this depends on the agent harness you are using, so ask your agent how to exclude certain files from being read by the agent.

Very often agents will use the default `.gitignore` file to exclude files from being read. But many have their custom special files such as `.ignore` (OpenCode), `.geminiignore` (Gemini CLI), `Codex does not have such special file yet`, read the documentation of your agent harness to see if it has such special file and how to use it.

But this is not a guarantee that the agent will not read these files, as there are many hacks that a rouge agent might use to bypass these ignore files. Also remember, that agents generally (at least by default) have write permissions to their own settings files, so there is also no guarantee that the agent will not change its settings to remove the ignore files or edit them.

7.3 Permission Systems

Even inside a sandbox, we may not want the agent to have “God Mode.” Permission systems act as a gatekeeper for sensitive actions. They can also act as a safety net in case you cannot setup a proper sandbox (e.g. you don’t have sufficient permissions to set up a sandbox on your work computer).

- **Human-in-the-Loop:** For high-risk actions (like deleting a directory or pushing to GitHub), the agent should ask for your permission first.
 - **Rule-Based Blocking:** Harnesses can be configured to automatically block access to certain folders or commands. For example, an agent might be allowed to edit your code but forbidden from reading your `.ssh` keys or `.env` files with passwords (e.g. credentials you may be using to access external services or databases). Also you might configure the agent harness to never execute file deletion commands, or commands that result in git commits or your project data upload to a remote git repository or other backup location (which may lead to data loss or corruption if agent messed something up locally).
-

7.4 Bottom line

It is very likely that even most sandboxed isolated agent customized with strict permissions will make a mistake at some point and do something you did not want it to do, such as deleting a file or installing a package that breaks your project. This is just the nature of current generation of agents, they are not perfect and they will make mistakes. So always make sure to have proper backups of your work, and if possible use version control (e.g. Git) so you can easily revert any unwanted changes.

For git, consider using authentication that requires user interaction for any changes done in your project locally to be pushed to your remote git repository or backup. Such precautions could be a password required for every access to ssh key for sending the data to the remote repository, or a biometric authentication (supported on some laptops and operating systems).

8. Task 2: Boundaries and Permissions

In this task, we will explore how to try controlling the agent's access to files and commands.

Open the `userspace/projects/task-02` folder in VS Code and open a terminal there. Start the agent harness there.

Prompt it:

```
I want to use HMDHFDplus R package to get the data for Germany for last few years.
I have password and user name, but i want the code in the R script that uses the
`HMDHFDplus` package to get the credentials from `.env` file. Can you write such script
and save it as `get_data.R`? I will put in my credentials int the `.env` file later,
for now just setup the code and that file with how you would expect the credentials
to be stored there. Double check list of functions in the pacakge and also using `?
<function_name>` check how to use the function that fetches the data.
```

Once this is set, read the documenation of the agent harness (e.g. for OpenCode this there is [a specific section with exact example on how to set this up](#), for Gemini CLI, just try googling 'in gemini cli how to make a hook that prevents if from reading .env file?'). Then pass over the instructions to the agent itself and ask it to configure itself so that it cannot read the `.env` file and thus cannot access the credentials stored there.

Once that is done, restart the agent.

Try to force it to read the file with natural language instructions. See if it works.

Try fooling it, by askig it to write an R script that reads any text file and prints it to console. Then manually edit this script to read the `.env` file and print it to console. Then ask the agent to run this script and tell you what the output is. See if it can read the credentials this way.

After all of these experiments, don't forget to change your password for HMD in case you did put in your real credentials into the `.env` file at some point.

III

Capabilities

9	Capabilities: Tools and Skills	41
9.1	Tools: The Agent's Hands	41
9.2	Skills: Specialized Capabilities	41
10	Task 2: Working with Tools and Skills	43
10.1	Setting the Scene	43
10.2	Step 1: Exploratory Analysis	43
10.3	Step 2: Activating a Skill	43
10.4	Step 3: Fixing a Bug	43

9. Capabilities: Tools and Skills

While the “Brain” (Model) can think and reason, it needs a way to interact with the world to be useful for coding, research or other tasks. These interaction is achieved through **Tools**, usually provided by the agent harness. **Skills** are higher-level bundles of instructions and logic that provide the agent with expertise in a specific domain, they may provide a specific way to do something or also add a custom tool (via shell, Python or any other script.)

9.1 Tools: The Agent’s Hands

Tools are discrete functions that the agent can execute to perform actions in your environment. Common tools include:

- **read_file**: Looking at the contents of a script or data file.
- **write_file / replace**: Creating new files or editing existing code.
- **run_shell_command**: Executing commands in the terminal (e.g., running an R script, installing a package, or listing files with `ls`, but also moving and deleting files, and virtually anything else you can do in the terminal - which is essentially anything).
- **web_search**: Accessing the internet to find documentation or troubleshooting tips.

The agent doesn’t just “run” these tools; it *reasons* about when to use them. For example, if you ask it to fix a bug, it will first use `read_file` to see the code, then `run_shell_command` to run a test that will produce output with the error message, it will think or search online or in the documentation how to fix and finally `replace` to attempt the fix. Then it is likely to run the test again to see if the fix worked, and if not it will repeat the process until the bug is fixed or it runs out of ideas or model quota, or, sometimes, model context window.

9.2 Skills: Specialized Capabilities

Skills are higher-level bundles of instructions and logic that provide the agent with expertise in a specific domain. While **tools** are the discrete actions an agent can perform, **skills** represent the programmed logic and best practices the agent follows to complete complex tasks potentially combining many standard tools or even custom scripts. This is very useful for cases when there is a very particular way something needs to be done, and the way you need the agent to do it is rather different how a general-purpose model would do it by following very frequent patterns in its training data.

Technically, **Skills** are just plain text files with metadata and textual instructions. Optionally there can also be scripts to be run during application of the skill.

For R language users, **Posit** (the maker of **RStudio**) has developed a set of skills and published them on GitHub at [posit/skills](https://github.com/posit-dev/skills). These skills are designed to help the agent understand how to perform common R programming tasks in a way that follows best practices and modern conventions. Follow the link and view any of the skills to see how they are structured and what kind of instructions they provide to the agent.

Skills can be global (stored in your user profile in a special folder) or local (stored in the current project folder). The agent automatically loads all global skills for all projects, and all project-specific skills when you are starting the agent in that project’s directory.

Tip

Activating Skills: You can explicitly ask an agent to activate a skill (e.g., “*Activate the R data analysis skill*”) to ensure it uses the most modern and idiomatic patterns for that task.

Skills should be concise, as they are loaded into the model’s context window. To save space in the context window, the agent only loads short metadata descriptions of each skill, and only loads the full instructions

when the agent decides it needs to use that skill. This allows you to have a large library of skills without overwhelming the model's context.

i Note

Note however, that having hundreds of skills might be counterproductive, as the agent might spend too much time evaluating which skills to load and use for a given task. It's best to have a curated set of skills that are relevant to your typical workflows or use project specific skills for very particular tasks that require a specific way of doing things.

💡 Tip

Anthropic (the maker of Claude) has initially introduced the concept of **Skills** and they have a special **Skill Creator** skill

10. Task 2: Working with Tools and Skills

In this task, you will use the Gemini CLI agent to explore a project, run code, and apply your first automated fix. We will use the `task-03-demotools` project.

10.1 Setting the Scene

Navigate to the project folder in your terminal:

```
cd userspace/projects/task-03-demotools
```

This folder contains a small R analysis and some sample data.

10.2 Step 1: Exploratory Analysis

Launch the agent and ask it to describe the project.

Prompt to try: > “Explore this directory. What data is here, and what does the `analysis.R` script do?”

Observation: Watch the agent use `ls` and `read_file` to build its understanding. It should identify the CSV file and the R script’s purpose.

10.3 Step 2: Activating a Skill

Let’s ensure the agent is in the right “mindset” for R data analysis.

Prompt to try: > “Activate the `r-data-analysis` skill and run the `analysis.R` script. Does it produce any output or errors?”

Observation: The agent will likely use `run_shell_command` to execute the R script. It might encounter an error if a package is missing or if there’s a bug in the code.

10.4 Step 3: Fixing a Bug

If the script failed (or even if it succeeded), let’s ask the agent to improve it.

Prompt to try: > “The script needs to calculate the growth rate of the population by region. Can you modify `analysis.R` to add this calculation and save the result to a new file called `growth_rates.csv`?”

Action: Watch the agent: 1. Read the current script. 2. Formulate a plan to add the `dplyr` logic. 3. Use the `replace` or `write_file` tool to update the code. 4. Run the script again to verify the output exists.

Tip: You can check the new `growth_rates.csv` file in your VS Code file pane to verify the agent’s work!

IV State Management

11	State Management: Context and Memory	47
11.1	The Context Window	47
11.2	Memory Systems	47
11.3	Agentic vs. Naive Approaches	47
12	Task 3: The Context Limit Challenge	49
12.1	Step 1: Generating the “Mega-File”	49
12.2	Step 2: The Naive Reality Check	49
12.3	Step 3: The Agentic Solution	49
12.4	Why this matters	49
12.4	References	51
12.4	Production tools	51
12.4	Website libraries	51
12.4	Local project code	52
12.4	Content license	52

11. State Management: Context and Memory

One of the biggest hurdles in working with LLMs is the **Context Window**. This is the limit on how much information (text, code, data) the model can process and reason over in a single pass.

11.1 The Context Window

Think of the context window as the agent's "working memory." - **Tokens**: Models don't read words; they read "tokens" (chunks of characters). A context window might be 32,000 tokens or 2 million tokens depending on the model. Check out the [OpenAI Tokenizer](#) to see how text is converted. - **Context Bloat**: If you try to paste an entire codebase or a massive dataset into a prompt, you will hit the context limit. Even if you don't hit the limit, very large contexts can make the model slower, more expensive, and less accurate (it might miss details in the middle).

11.2 Memory Systems

To handle large projects, agents use **Memory Systems**. Instead of loading everything into the "working memory" at once, they: - **Index** your files (creating a searchable map of your project). - **Search** for relevant information only when needed. - **Store** past decisions and discoveries in a long-term memory file (like `MEMORY.md` or a local database) so they don't have to re-learn things in the next session.

11.3 Agentic vs. Naive Approaches

When faced with a massive file: - **Naive User**: Copies and pastes the whole file into ChatGPT. Result: "Message too long" or truncated response. - **Agentic Tool**: Uses tools like `grep_search` to find specific keywords or `read_file` with `start_line` and `end_line` to read only the relevant section.

Note

Think-in-Code: Modern agents are taught to "think in code." Instead of reading 100MB of logs, they write a small script to summarize the logs and only read the summary. This keeps the context window clean for actual reasoning.

In the next exercise, we will purposely try to "break" the context window using a massive documentation dump and see how the agent handles it.

12. Task 3: The Context Limit Challenge

In this task, we will simulate a real-world scenario where you have “too much information” for the model to handle at once. We’ll use the `rdocdump` tool to create a massive text file from an R package.

12.1 Step 1: Generating the “Mega-File”

Open your terminal and run the following command to dump the entire documentation of the `DemoTools` package into a single text file:

```
# This uses rdocdump to pull the full text of the package from GitHub
# Note: You might need to install it first if not in the container
# remotes::install_github("e-kotov/rdocdump")
Rscript -e 'rdocdump::rdocdump("e-kotov/DemoTools", out = "demotools_dump.txt")'
```

12.2 Step 2: The Naive Reality Check

Now that you have `demotools_dump.txt`, try the “Naive” approach:

1. Open the file in VS Code and copy **all** the text (it will be thousands of lines).
2. Go to the [OpenAI Tokenizer](#) and paste it. How many tokens is it?
3. Try pasting it into [Google AI Studio](#) or ChatGPT.
 - Does it fit?
 - Is it hard to scroll?
 - Imagine doing this for every question you have!

12.3 Step 3: The Agentic Solution

Now, let’s see how an agent handles this without blowing up the context window.

Launch the Gemini CLI or OpenCode agent and give it a broad task.

Prompt to try: > “I have a file called `demotools_dump.txt`. Read it and explain how the `DemoTools` package handles life table calculations. Provide a code example based on what you find.”

Observation: The agent **cannot** read the entire file at once because it has a safety limit (usually around 2000 lines). Watch how it pivots: - It might use `grep_search` to find “life table”. - It might read the file in small chunks using `start_line` and `end_line`. - It might use a “Context Mode” tool to index the file in a sandbox and only return the relevant sections.

12.4 Why this matters

By using an agent, you avoided: 1. **Wasting money/tokens** on thousands of lines of irrelevant documentation. 2. **Confusing the model** with too much noise. 3. **Manual effort** of hunting through a giant text file.

Tip: If the agent gets stuck, help it! You can say: “Use `grep_search` to find the life table functions in the dump file.”

References

This page lists the direct software, libraries, and local helper scripts used to build and display this book. It is intended as a practical attribution and reproducibility note, not a full software bill of materials for every transitive dependency bundled by each tool.

Production tools

Software	Version used locally	Role	License / project
Quarto	1.9.37	Builds the book website and PDF from the <code>.qmd</code> source files.	Open-source CLI Quarto
Pandoc	3.8.3	Document conversion engine used by Quarto.	GPL-2.0-or-later
Typst	0.14.2	PDF typesetting backend used by the Quarto <code>typst</code> format.	Apache-2.0

Website libraries

Software	How it is used	License / project
Typed.js 2.1.0	Loaded from unpkg.com for the animated “Beyond the Chatbox” title on the home page.	MIT
Bootstrap and the Cosmo theme	HTML layout, responsive navigation, typography, and theme styling through Quarto.	MIT
Bootstrap Icons	Icons used in Quarto navigation and controls.	MIT
Clipboard.js	Code-copy buttons generated by Quarto.	MIT
Fuse.js and autocomplete support	Client-side search generated by Quarto.	Apache-2.0
Tippy.js and Popper	Tooltips and popovers generated by Quarto.	MIT , MIT
AnchorJS	Header anchor links generated by Quarto.	MIT
Headroom.js	Navigation header behavior generated by Quarto.	MIT

i Typed.js license version

This book intentionally pins Typed.js to version 2.1.0, whose release tag includes an MIT license. The current `main` branch of Typed.js uses a GPL-3.0-or-later license. Re-check the upstream license before upgrading the CDN URL to a newer Typed.js release.

Local project code

The book also uses small local scripts and styles maintained in this repository:

File	Purpose
<code>styles.css</code>	Custom visual styling, including the Typed.js cursor styling.
<code>scripts/typed-init.js</code>	Initializes the Typed.js title animation.
<code>scripts/details-fold.js</code>	Adds behavior for collapsible details blocks.
<code>scripts/trigger-fix.lua</code> and <code>scripts/fix-typst-bg.lua</code>	Support the PDF/Typst rendering workflow.
<code>scripts/analytics.html</code>	Adds site analytics markup.

Content license

Unless otherwise noted, the workshop materials are distributed under [CC-BY-SA 4.0](#), as stated in the site footer. Third-party software, images and other content remains under their respective upstream licenses.