

Beyond the Chatbox

Egor Kotov

Contents

	Introduction	5
	Overview	5
	Workshop Details	5
	Before the Workshop: Accounts & Setup	6
	1. Recommended Accounts to Register	6
	2. Coding Agents We Will Use (and Why)	6
1	Introduction to LLM Agents	9
1.1	What is an LLM Agent?	9
1.1.1	Key Characteristics	9
1.2	Why Use Agents for Research?	9
1.3	Models	9
1.4	Next Steps	9
2	Security & Sandboxing	11
2.1	Common Risks	11
2.1.1	Unsafe Code Execution	11
2.1.2	Data Leakage	11
2.1.3	Environment Pollution	11
2.2	Mitigation Strategies	11
2.2.1	Containerization (Docker)	11
2.2.2	Managed Environments (GitHub Codespaces)	11
2.2.3	The “Human in the Loop”	11
2.3	Practice Safe Coding	11
3	Environment Setup	13
3.1	 Launching the Environment	13
3.1.1	 Auto-Pause & Resume	14
3.2	 Managing Your Session & Quotas	14
3.2.1	Monthly Free Quotas (2026)	14
3.2.2	Checking Remaining Quota	14
3.2.3	Stopping Your Session Manually	14
3.3	Next Steps	15
4	The Workshop Workspace	17
4.1	Workspace Structure	17
4.2	The Sandbox Boundary	17
4.3	Next Steps	17

5	Interacting with the Agents	19
5.1	Available CLI Agents	19
5.2	Exercise 1: Exploring Data with Gemini	19
5.3	Exercise 2: Pair Programming with Aider	19
5.4	Exercise 3: Rapid Prototyping with NCA	19
5.5	Security & Best Practices Reminder	19
5.6	Next Steps	19

Introduction

⚠ Warning

THIS IS STILL WORK IN PROGRESS - EXPECT BROKEN LINKS AND MISSING CONTENT



Figure 0.1: © Image generated by Google Gemini | Nanobanana

Overview

This interactive tutorial and workshop provides a practical introduction to using large language models (LLMs) in coding agents, featuring applied examples in population research. It is organized by [Egor Kotov](#) with support from the [Max Planck Institute for Demographic Research](#).

Workshop Details

- **Date:** June 03, 2026
- **Time:** 13:30-16:30
- **Location:** Complesso Belmeloro, room “Lab O”, [Via Beniamino Andreatta, 8, 40126 Bologna BO, Italy](#)

Abstract

The software industry has been enjoying the benefits of coding agents for over a year now, with even top software engineers now having AI assistants generate a vast majority of their code. Now, the tools for agentic coding are mature enough for anyone to try.

Unlike many educational materials in this field, this workshop will also address the new vectors of attack that novice users might expose themselves to by using coding agents blindly. We will cover key security risks associated with agentic systems, including data leakage, unsafe code execution, and unintended modifications to analysis code and research materials. To mitigate these risks, we will introduce accessible strategies such as containerization, sandboxing, and workflow isolation.

Participants will also see how the agents themselves can assist in setting up and using these protective tools to improve transparency and reproducibility. The workshop will combine short demonstrations with applied examples relevant to population research. Participants are encouraged to bring their own projects or ideas, which we will use to explore how LLM agents can be integrated into real research workflows. For those without a project, guided exercises and example tasks will be provided.

Before the Workshop: Accounts & Setup

To make the most of our hands-on sessions, please complete the following steps before the workshop on **June 3, 2026**.

! Data Privacy & Security Warning

Free tiers for these LLM services typically **use your prompts and uploaded code for model training**. **Do not use these tools on private, sensitive, or proprietary research data!** Only use public or non-sensitive datasets during the workshop and for general testing.

1. Recommended Accounts to Register

Please sign up for the following services in advance to ensure you have enough free model usage quota during the exercises:

- **Google Account** (A personal/secondary account is perfect) — Used for **Gemini CLI** and **Antigravity CLI**. No credit card is required.
- **OpenCode Zen** (opencode.ai/zen) — Gives you free model access via the **OpenCode** agent. You can view the list of available models and details on [OpenCode Zen pricing/documentation](#).
- **OpenRouter** (openrouter.ai) — Allows accessing a wide variety of models. You'll need to generate a free **API Key** after registering, which can be used to access their catalog of **free models**.
- *Optional: NVIDIA NIM/Build* (build.nvidia.com) — Offers **free preview models** (note: account activation requires SMS verification via a mobile phone number).

i Note

If you have a subscription to another LLM service (such as OpenAI ChatGPT Plus, Claude Pro, Gemini Advanced, or Mistral) and prefer to use it during the workshop, please feel free to do so. However, our guided exercises will focus on the tools listed above to ensure everyone has access to the same models and features.

The skills you will learn here are highly transferable to other AI coding assistants, though setup steps and specific features may vary by provider. Please refer to the documentation of your chosen agent for setup instructions, and don't hesitate to ask the agent itself how to configure it for your desired models and features.

⚠ API & Account Safety Warning

Do **NOT** attempt to authenticate your Google or Anthropic (Claude) accounts directly within OpenCode. Their terms of service prohibit use of quota based subscriptions in third-party clients, and doing so can result in your accounts being permanently banned. OpenAI is currently a bit more flexible on this at the time of writing, but always verify provider policies.

2. Coding Agents We Will Use (and Why)

The coding agents are already pre-installed and configured inside the workshop's Docker container / Codespaces environment. We will focus on:

- **Google Antigravity CLI** (and its predecessor **Gemini CLI**)
 - *Why:* It provides a free daily quota to test cutting-edge models (like **Gemini 3.5 Flash**, **Gemini 3.1 Pro**, and **Claude 3.5 Sonnet**) directly from your Google account.

- **OpenCode CLI**

- *Why:* A flexible, open-source-friendly agent that can connect to multiple backends. We will use it with free models from **OpenCode Zen**, **OpenRouter**, and **NVIDIA NIM** to show how easy it is to switch between different model providers without vendor lock-in.

i Note

What is **CLI**? It stands for “Command Line Interface” and means that you will interact with the agent by typing commands in a terminal (similar to your R console), rather than through a graphical user interface (GUI). This allows for more flexibility and control over the agent’s behavior, as well as easier integration into coding workflows. This also gives you better insight into how the agent works under the hood, which is important for understanding its capabilities and limitations. You are of course free to use the same agents through their graphical interfaces too. We are also limited by the current tutorial environment and tools that are free to use. In our tutorial environment Gemini and Antigravity are only available as CLI, OpenCode is also CLI, but you will learn how to run it in graphical mode as well.

To get started with the workshop materials, please navigate through the sections in the navigation bar on the left or with pagination links in the bottom.

1. Introduction to LLM Agents

This section covers the basics of what LLM agents are and how they differ from standard chat interfaces.

1.1 What is an LLM Agent?

Unlike a standard LLM chat (like ChatGPT or Gemini web interface), an **agent** is a system that uses an LLM as its “reasoning engine” to perform tasks by interacting with the world.

1.1.1 Key Characteristics

- **Tool Use:** Agents can use tools such as search engines, calculators, or terminal commands.
- **Autonomy:** They can break down a complex goal into smaller steps and execute them sequentially.
- **Feedback Loop:** They can observe the output of a tool (e.g., a compiler error) and adjust their next action accordingly.

1.2 Why Use Agents for Research?

In academic research, agents can: - **Automate Data Cleaning:** Write and run scripts to fix inconsistencies in large datasets. - **Literature Review:** Search, summarize, and synthesize findings from hundreds of papers. - **Reproducible Analysis:** Generate documented code that follows best practices.

1.3 Models

See [models comparison](#)

1.4 Next Steps

Now that you know what agents are, let's look at [Security & Sandboxing](#).

2. Security & Sandboxing

Using coding agents involves granting them the ability to execute code on your machine. This introduces several security and integrity risks that must be managed.

2.1 Common Risks

2.1.1 Unsafe Code Execution

An agent might generate code that deletes files, installs malicious packages, or opens network vulnerabilities, especially if it “hallucinates” a solution or misinterprets a command.

2.1.2 Data Leakage

When you provide an agent with access to your files, it might send sensitive data (API keys, personal identifiers, unpublished research) to the LLM provider’s servers for processing.

2.1.3 Environment Pollution

Agents might install many dependencies or modify system configurations, making your research environment unstable or non-reproducible.

2.2 Mitigation Strategies

2.2.1 Containerization (Docker)

By running the agent inside a container, you limit its access to your host machine’s files and resources.

2.2.2 Managed Environments (GitHub Codespaces)

Using a temporary, cloud-based environment like Codespaces provides a “disposable” workspace where you can safely experiment without risking your local setup.

2.2.3 The “Human in the Loop”

Never allow an agent to execute code autonomously without a review step. Tools like **Aider** or **Gemini CLI** often prompt for confirmation before running critical commands.

2.3 Practice Safe Coding

Our [Environment Setup](#) is hosted entirely within a secure GitHub Codespace to ensure your experiments are isolated. Let’s head there now to start exploring.

3. Environment Setup

To ensure a consistent, secure, and hassle-free experience, we will use a pre-configured cloud-based development environment. We will use **GitHub Codespaces**.

💡 ⏳ Start the Setup Early and Focus on the Presentations!

First-time environment setup and container building can take **up to 10 minutes** to compile and prepare.

We highly recommend launching the codespace right at the beginning of the workshop. While the setup runs automatically in the background, you can continue to fully focus on the presentations, live demonstrations, and slides shown by the tutor.

3.1 🛠️ Launching the Environment

To launch the Codespaces environment manually from the GitHub interface, follow these steps:

1. Navigate to the [workshop repository](#).
2. Click the green **Code** button in the top right.
3. Select the **Codespaces** tab.
4. Click **Create codespace on main** (or click the **...** menu to configure machine size).




Figure 3.2: Launch in GitHub Codespaces

Free & Secure Environment

- **It is 100% Free:** GitHub provides a generous free tier of **60 hours** of Codespaces usage per month for every personal account. No credit card, billing info, or payment is required.
- **It is Completely Secure:** The environment runs inside a secure, isolated container (virtual machine) in the cloud. It does not touch your physical computer, install any local files, or run any software on your machine. It belongs to you and it is also mostly safe to enter your account data (e.g. Google account, especially if secondary or test account, API keys, other accounts for subscription services or services we will use during the workshop) without worrying about security or privacy issues.

3.1.1 Auto-Pause & Resume

To save your free hours, the environment automatically pauses after 30 minutes of inactivity. When you restart it, your files, open tabs, and terminal history are fully preserved—just like waking up a laptop from sleep (though you may need to restart any active terminal commands or running applications).

3.2 Managing Your Session & Quotas

GitHub Codespaces provides a generous free tier, but it is important to understand how usage is consumed to avoid running out of credits.

3.2.1 Monthly Free Quotas (2026)

The following quotas apply to personal accounts.¹

Account Type	Actual Usage Hours	Storage (per month)
GitHub Free	60 hours	15 GB
GitHub Pro	90 hours	20 GB

Storage Persistence Quota

Storage (15GB/20GB) is consumed even when the codespace is **stopped**. To save your storage quota, make sure to **delete** codespaces you no longer need at github.com/codespaces once the workshop is fully over.

3.2.2 Checking Remaining Quota

Tip

You can monitor your real-time usage and remaining balance at any time:

1. Go to your [GitHub Settings > Billing and plans](#).
2. Look for the **Codespaces** section under **Usage this month** or try this direct link: [GitHub Codespaces Usage](#).

3.2.3 Stopping Your Session Manually

When you are done with a session, you can stop the codespace manually to instantly stop consuming usage hours:

¹While the [official billing documentation](#) lists these as 120 and 180 **core hours**, the minimum instance size is 2 CPU cores. This means every 1 hour of real-time work consumes 2 core hours, resulting in the actual limits shown below.




Figure 3.3: Stop GitHub Codespace

3.3 Next Steps

Now that your environment is ready, let's explore the [Workshop Workspace](#).

4. The Workshop Workspace

Once your environment is ready, you will find a dedicated space for your work. Understanding the structure of this workspace is crucial for maintaining security and reproducibility.

4.1 Workspace Structure

The root directory contains several key folders:

- `userspace/`: This is your personal sandbox.
 - `userspace/examples/`: contains pre-built examples that we will explore together.
 - `userspace/projects/`: This is where you can start your own research projects or experiments.
- `data/`: Contains shared datasets for the workshop exercises.
- `tutorial-website-src/`: The source code for this website (if you want to see how it's built!).

4.2 The Sandbox Boundary

As we discussed in [Security & Sandboxing](#), the `userspace/` directory is designed to be your primary work area.

! Best Practice

When instructing an agent to create or modify files, always specify a path within `userspace/projects/`. This keeps your experimental code isolated from the core workshop materials.

4.3 Next Steps

Now that you understand where to work, let's learn how to [Interact with the Agents](#).

5. Interacting with the Agents

In this environment, you have access to several world-class coding agents. You can interact with them via the terminal or integrated VS Code extensions.

5.1 Available CLI Agents

- **Aider:** A powerful pair-programming tool that works directly with your git repository. Run `aider` in the terminal.
- **Gemini CLI:** Interactive chat with Google's Gemini models. Run `gemini-chat`.
- **Claude CLI:** Direct access to Anthropic's Claude models. Run `claude`.
- **NCA (Nim Coding Agent):** Optimized for speed and specific tasks. Run `nca`.

5.2 Exercise 1: Exploring Data with Gemini

Try asking Gemini to help you understand the shared datasets. Open your terminal and run:

```
gemini-chat "List the files in the data/ directory and give me a brief summary of the columns in the first CSV file you find."
```

5.3 Exercise 2: Pair Programming with Aider

Aider is excellent for editing existing code or starting new scripts. Let's try creating a script in your `userspace/projects/` folder:

1. Open the terminal.
2. Run `aider userspace/projects/analysis.py`.
3. Inside the Aider chat, type: `> "Read data/sample_data.csv and create a plot of the population trends by year."`

5.4 Exercise 3: Rapid Prototyping with NCA

If you need a quick script or a one-liner, `nca` is very efficient:

```
nca "Write a python one-liner to count the number of rows in data/sample_data.csv"
```

5.5 Security & Best Practices Reminder

As discussed in [Security & Sandboxing](#), always remember:

- **Work within userspace/:** This ensures your experiments are isolated.
- **Review Agent Code:** Never execute code generated by an agent without reviewing it first.
- **Protect Sensitive Data:** Do not provide agents with access to private credentials or sensitive personal data.

5.6 Next Steps

Now that you are comfortable with the agents, let's move on to the [guided examples](#) (coming soon).